



Advanced UNIX System Programming

2

Agenda

1 Daemon Process



2 System V IPC



3 Socket Programming



Daemon Process

Introduction

- Daemon process starts during system startup
- They frequently spawn other process to handle services requests
 - Mostly started by initialization script /etc/rc
- Waits for an event to occur
- perform some specified task on periodic basis (cron job)
- perform the requested service and wait
 - Example print server

Characteristics

- executed at the background process
- Orphan process
- No controlling terminal
- run with super user privileges
- process group leaders
- session leaders

Daemon -program

```
int init_daemon ( void ) {  
    if ( ! fork ( ) ) {  
        setsid ( ) ;  
        chdir ( " / " ) ;  
        umask ( 0 ) ;  
        /* Specify Your Job */  
        return ( 0 ) ;  
    }  
    else  
        exit ( 0 ) ;  
}
```

System V IPC

Introduction

- Pipe and FIFO do not satisfy many requirements of many applications.
- Sys V IPC is implemented as a single unit
- System V IPC Provides three mechanisms namely
 - Message Queues,
 - Shared Memory
 - Semaphores.
- Persist till explicitly delete or reboot the system

Common Attributes

- Each IPC objects has the following attributes.
 - key
 - id
 - Owner
 - Permission
 - Size
 - Message queue – used-bytes, number of messages
 - Shared memory – size, number of attach, status
 - Semaphore – number of semaphores in a set
- The ipc_perm structure holds the common attributes of the resources.

System Limitations

```
$ ipcs -l
```

```
----- Shared Memory Limits -----
```

```
max number of segments = 4096
```

```
max seg size (kbytes) = 32768
```

```
max total shared memory (kbytes) = 8388608
```

```
min seg size (bytes) = 1
```

```
----- Semaphore Limits -----
```

```
max number of arrays = 128
```

```
max semaphores per array = 250
```

```
max semaphores system wide = 32000
```

```
max ops per semop call = 32
```

```
semaphore max value = 32767
```

```
----- Messages: Limits -----
```

```
max queues system wide = 16
```

```
max size of message (bytes) = 8192
```

```
default max size of queue (bytes) = 16384
```

Get a key

- If we wish to communicate between different processes using an IPC resource, the first step is to create a shared unique identifier.
- The system generates a number dynamically for a given mechanism by using the *ftok* library function.
- But apart from the creator, other processes that want to communicate with the creator process should agree to the key value.
- Syntax: `key_t ftok (const char *filename, int id);`

Get an id

- The syntax for a *get* function is:
int xxxget (key_t key, int xxxflg);
(xxx may be msg or shm or sem)
- If successful, returns to an identifier; otherwise -1 for error.
- The key can be generated in three different ways
 - from the *ftok* library function
 - by choosing some static positive integer value
 - by using the IPC_PRIVATE macro
- flags commonly used with this function are IPC_CREAT and IPC_EXCL.

Control a object

- The syntax for the *control* function is:

```
int xxxctl (int xxxid, int cmd, struct xxxid_ds *buffer);  
(xxx may be msg or shm or sem);
```

- If successful, the *xxxctl* function returns zero, otherwise it returns -1.
- The command argument may be
 - IPC_STAT
 - IPC_SET
 - IPC_RMID

Message Queues - Introduction

- Message queue overcomes FIFO limitation like storing data and setting message boundaries.
- Create a message queue
- Send message (s) to the queue
- Any process who has permission to access the queue can retrieve message (s).
- remove the message queue.

Message queues

```
struct msgbuf {  
    long mtype;  
    char mtext [1];
```

}; Standard structure

```
struct My_msgQ {  
    long mtype;  
    char mtext [1024];  
    void * xyz;
```

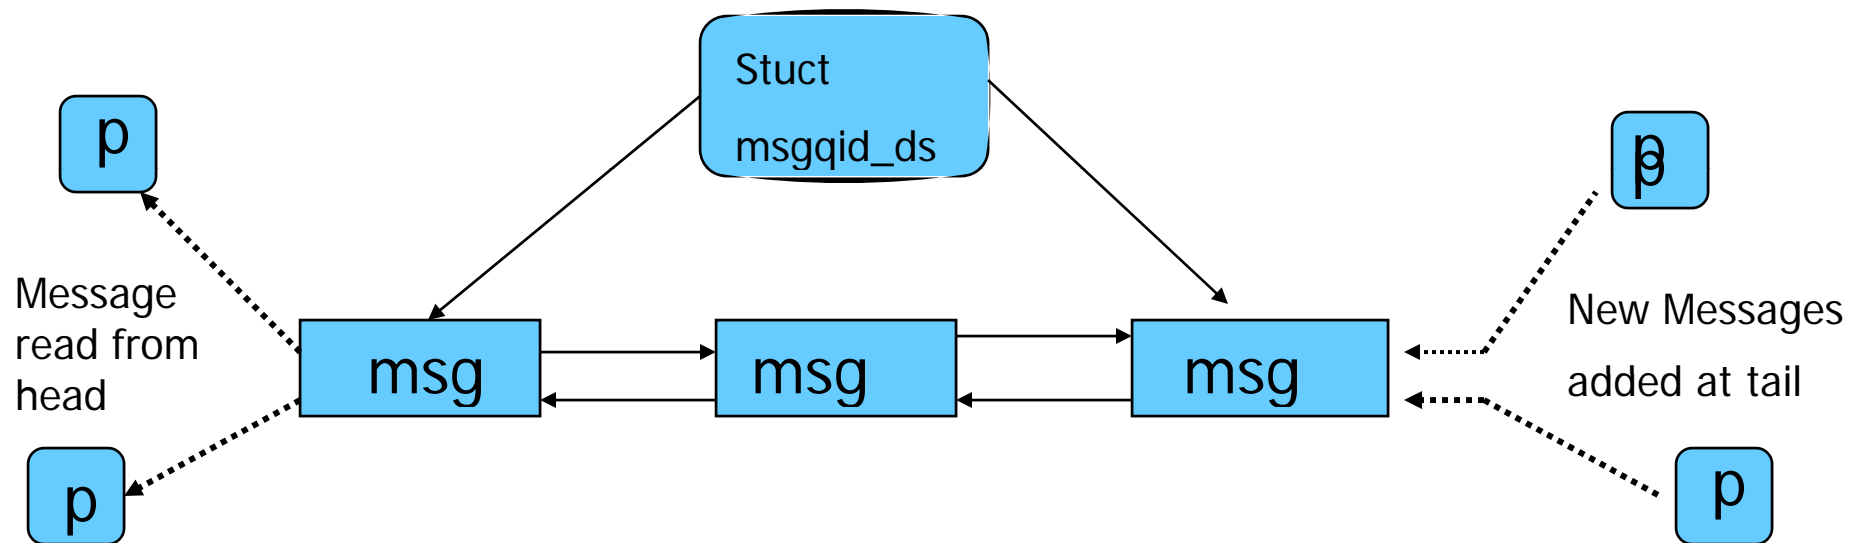
}; Our own structure

msqid xxx

mtype x_1	msg text
mtype x_2	msg text
mtype x_3	msg text
mtype x_4	msg text
mtype x_5	msg text

mtype x_n	msg text

Messages in a queue



msqid_ds

```
struct msqid_ds
{
    struct ipc_perm msg_perm;
    __time_t msg_stime;
    __time_t msg_rtime;
    __time_t msg_ctime;
    unsigned long int __msg_cbytes;
    msgqnum_t msg_qnum;
    msglen_t msg_qbytes;
    __pid_t msg_lspid;
    __pid_t msg_lrpid;
};
```

msgget

- **int msgget (key_t key, int msgflg);**
- The first argument key can be passed from the return value of the ftok function or made IPC_PRIVATE.
- To create a message queue, IPC_CREAT ORed with access permission is set for the msgflg argument.
- Ex: msgid = msgget (key, IPC_CREAT | 0744);
 msgid = msgget (key, 0);

msgsnd

- The syntax of the function is:

```
int msgsnd (int msqid, struct msgbuf *msgp, size_t msgsz, int msgflg);
```

- Arguments:

- message queue ID
- address of the structure.
- size of the message text
- message flag
 - 0 or IPC_NOWAIT

msgrcv

- syntax of the function is:

```
ssize_t msgrcv (int msqid, struct msgbuf *msgp, size_t msgsz, long msgtype,  
int msgflg);
```

- msgtype argument is used to retrieve a particular message.
 - 0 -retrieve in FIFO order
 - +ve - retrieve the the exact value of the message type
 - -ve - first message or <= to the absolute value.
- on success, msgrcv returns with the number of bytes actually copied into the message text

Destroying a message queue

- There are many ways:
- From command line, using one of the ways
 - \$ ipcrm msg msqid
 - \$ ipcrm -q msqid
 - \$ ipcrm -Q msgkey
- Using system call
 - msgctl (msgid, IPC_RMID, 0);

Message queue : pseudo code

- `key = ftok (".", 'a');`
- `msqid = msgget (key, IPC_CREAT|0666);`
- `msgsnd (msqid, &struct, sizeof (struct), 0);`

- `msgrcv (msqid, &struct, sizeof (struct), mtype, 0);`
- `msgctl (msqid, IPC_RMID, NULL);`
- `$ipcrm msg msqid`

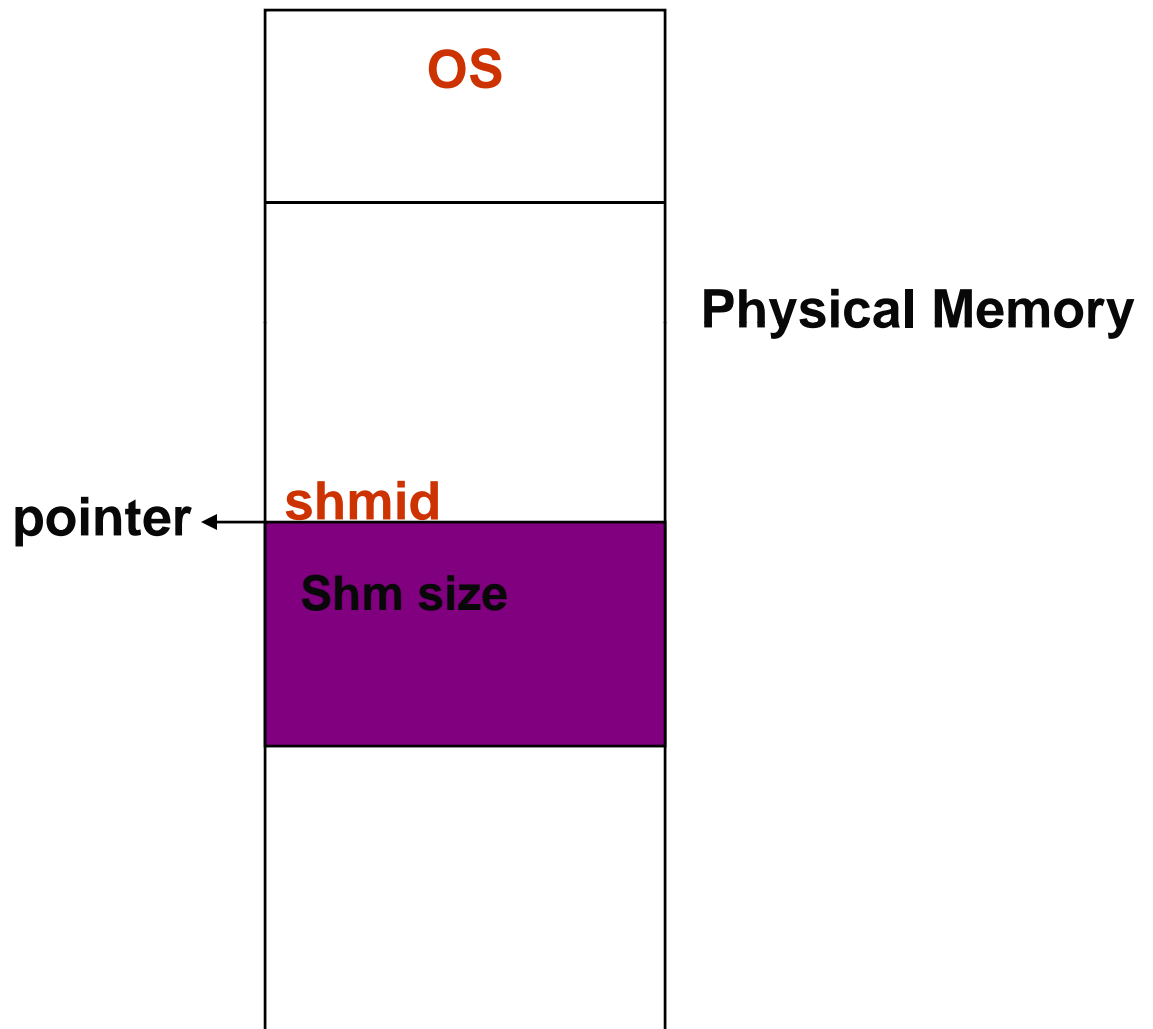
Limitations

- Message queues are effective if a small amount of data is transferred.
- Very expensive for large transfers.
- During message sending and receiving, the message is copied from user buffer into kernel buffer and vice versa
- So each message transfer involves two data copy operations, which results in poor performance of a system.
- A message in a queue can not be reused

Shared Memory - Introduction

- Very flexible and ease of use.
- Fastest IPC mechanisms
- shared memory is used to provide access to
 - Global variable
 - Shared libraries
 - Word processors
 - Multi-player gaming environment
 - Http daemons
 - Other programs written in languages like Perl, C etc.,

Shared Memory – shmat ()



Why go for Shared Memory

- Shared memory is a much faster method of communication than either semaphores or message queues.
- does not require an intermediate kernel buffer
- Using shared memory is quite easy. After a shared memory segment is set up, it is manipulated exactly like any other memory area.

Shared memory - Data structures

- The data structures used in shared memory are:
 - shmid_ds
 - ipc_perm
 - shminfo
 - shm_info
 - shmid_kernel

ipc_perm structure

```
struct ipc_perm
{
    __key_t __key;           - Key
    __uid_t uid;             - Owner's user ID
    __gid_t gid;             - Owner's group ID
    __uid_t cuid;            - Creator's user ID
    __gid_t cgid;            - Creator's group ID
    unsigned short int mode;  - r/w permission unsigned
    short int __seq;         - Sequence number
};
```

shmid_ds

```
struct shmid_ds
{
    struct ipc_perm shm_perm;
    size_t shm_segsz;
    __time_t shm_atime;
    __time_t shm_dtime;
    __time_t shm_ctime;
    __pid_t shm_cpid;
    __pid_t shm_lpid;
    shmatt_t shm_nattch;
};
```

Steps to access Shared Memory

- The steps involved are
 - Creating shared memory
 - Connecting to the memory & obtaining a pointer to the memory
 - Reading/Writing & changing access mode to the memory
 - Detaching from memory
 - Deleting the shared segment

shmat

- used to attach the created shared memory segment onto a process address space.
- `void *shmat(int shmid,void *shmaddr,int shmflg)`
- Example: `data=shmat(shmid,(void *)0,0);`
- A pointer is returned on the successful execution of the system call and the process can read or write to the segment using the pointer.

Reading/ writing to SM

- Reading or writing to a shared memory is the easiest part.
- The data is written on to the shared memory as we do it with normal memory using the pointers
- Eg. Read:
 - `printf("SHM contents : %s \n", data);`
- Write:
 - `printf("Enter a String : ");`
 - `scanf(" %[^\n]",data);`

Shmdt & shmctl

- The detachment of an attached shared memory segment is done by shmdt to pass the address of the pointer as an argument.
- Syntax: `int shmdt(void *shmaddr);`
- To remove shared memory call:
`int shmctl(shmid,IPC_RMID,NULL);`
- These functions return `-1` on error and `0` on successful execution.

Shared Memory – pseudo code

- `shmid = shmget (key, 1024, IPC_CREAT|0744);`
- `void *shmat (int shmid, void *shmaddr, int shmflg);`
if the shm is read only pass SHM_RDONLY else 0
- `(void *)data = shmat (shmid, (void *)0, 0);`
- `int shmdt (void *shmaddr);`
- `int shmctl (shmid, IPC_RMID, NULL);`

Limitations

- Data can either be read or written only. Append is not allowed
- Race condition
 - Since many processes can access the shared memory, any modification done by one process in the address space is visible to all other processes. Since the address space is a shared resource, the developer should implement a proper locking mechanism to prevent the race condition in the shared memory.

Semaphores

Synchronization Tool

An Integer Number

P () And V () Operators

Avoid Busy Waiting


Types of Semaphore

Used in :

shared memory segment

message queue

file



Semaphores (Contd.).

- If a process wants to use the shared object, it will “lock” it by asking the semaphore to decrement the counter
- Depending upon the current value of the counter, the semaphore will either be able to carry out this operation, or will have to wait until the operation becomes possible
- The current value of counter is >0 , the decrement operation will be possible. Otherwise, the process will have to wait

System V IPC – Semaphores

- System V semaphore provides a semaphore set - that can include a number of semaphores. It is up to user to decide the number of semaphores in the set
- Each semaphore in the set can be a binary or a counting semaphore. Each semaphore can be used to control access to one resource - by changing the value of semaphore count

Semaphore - Initialization

```
union semun {  
    int val;                // value for SETVAL  
    struct semid_ds *buf;    // buffer for IPC_STAT, IPC_SET  
    unsigned short int *array; // array for GETALL, SETALL  
};
```

```
union semun arg;
```

```
semid = semget (key, 1, IPC_CREAT | 0644);  
arg.val = 1; /* 1 for binary else > 1 for Counting Semaphore */  
semctl (semid, 0, SETVAL, arg);
```

Semaphore - Implementation

```
struct sembuf {  
    short sem_num; /* semaphore number: 0 means first */  
    short sem_op; /* semaphore operation: lock or unlock */  
    short sem_flg; /* operation flags : 0, SEM_UNDO, IPC_NOWAIT */  
};  
struct sembuf buf = {0, -1, 0}; /* (-1 + previous value) */  
semid = semget (key, 1, 0);  
  
semop (semid, &buf, 1); /* locked */  
-----Critical section-----  
buf.sem_op = 1;  
semop (semid, &buf, 1); /* unlocked */
```

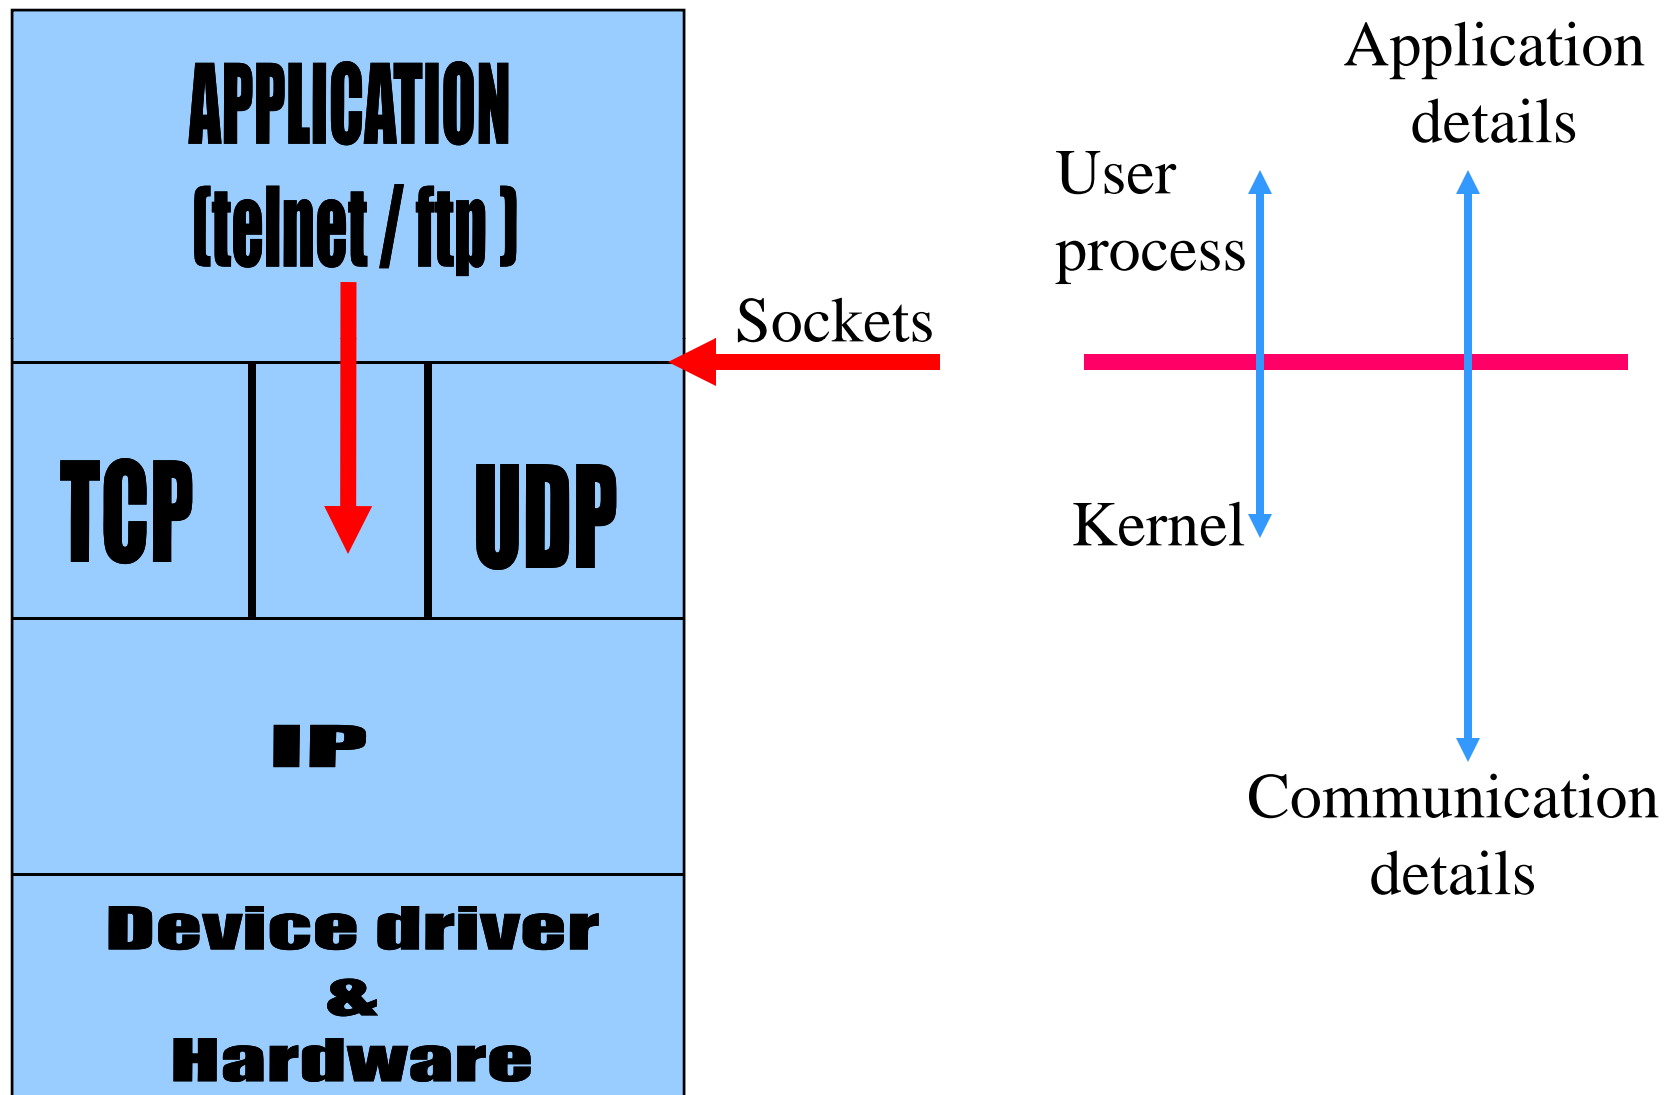

Socket Programming

Socket


A socket is used to communicate between different machines (different IP addresses). Socket of type SOCK_STREAM is full-duplex byte streams.



TCP/ IP Protocol Stack



Concurrent Server



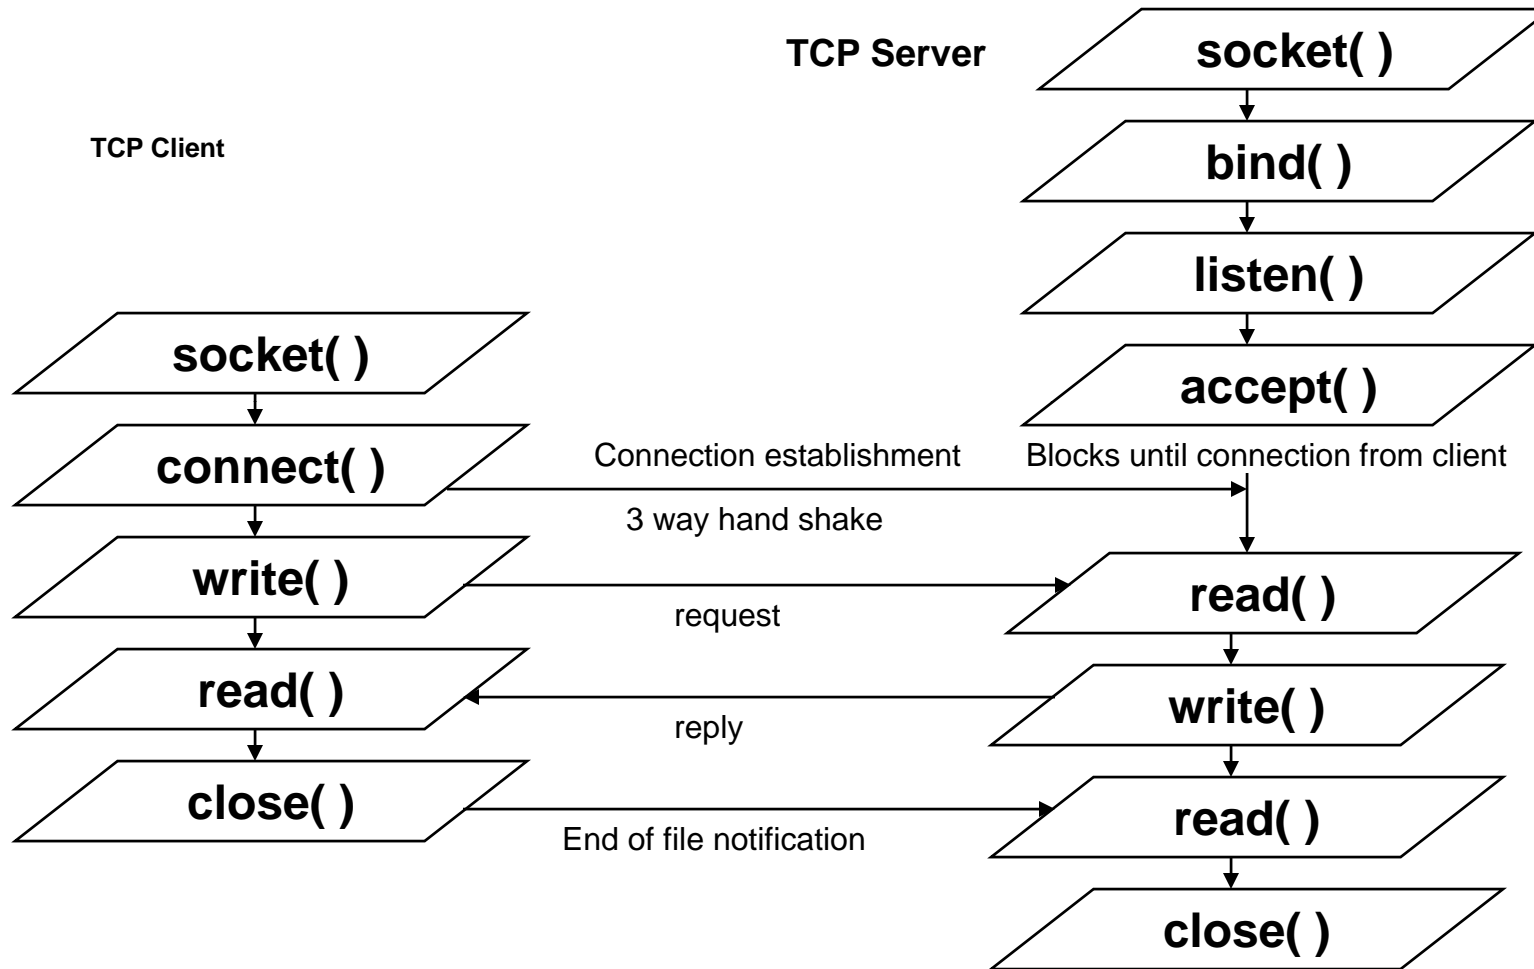
Sockets (Contd.).

- A socket is a communication endpoint and represents abstract object that a process may use to send or receive messages.
- The socket frame work is procedural not message based
- The two most prevalent communication APIs for Unix Systems are Berkeley Sockets and System V Transport Layer Interface(TLI)

Sockets (Contd.).

- The typical client -server relationship is not symmetrical
- Network Connection can be connection-oriented or connectionless
- More parameters must be specified for network connection, than for file I/O
- The unix I/O system is stream oriented
- The network interface should support multiple communication protocol

Socket Functions



socket ()

- `int socket (int domain, int type,
 int protocol);`
- Domain
 - `AF_UNIX` for UNIX domain
 - `AF_INET` for Internet domain
- Socket type
 - `SOCK_STREAM` for TCP (Connection Oriented)
 - `SOCK_DGRAM` for UDP (Connectionless)
- Protocol
 - Protocol number is used to identify an application. List of the protocol number and the corresponding applications can be seen at `/etc/protocols`.
- The socket system call returns a socket descriptor on success and -1 for failure.


sockaddr_in structure

- ```
struct sockaddr_in {
 short int sin_family;
 unsigned short int sin_port;
 struct in_addr sin_addr;
}
```
- sin\_family - address family
- sin\_port - port number
- sin\_addr - internet address (IP addr)
  
- The in\_addr structure used to define sin\_addr is as under  


```
struct in_addr {
 unsigned long s_addr; /* refers to the four byte IP address */
}
```

Byte ordering ex: 91,329 hex: 00 01 64 C1

**Little endian byte order :**    **Ex: Intel series**



**Big endian byte order :**    **ex: IBM 370, Motorola**



# Byte ordering functions

---

- Internet protocols use big endian byte ordering called network byte order
- Functions below allow conversions between the formats.

```
#include <netinet/in.h>
```

htons() – "Host to Network Short"

htonl() – "Host to Network Long"

ntohs() – "Network to Host Short"

ntohl() – "Network to Host Long"

- h stands for host      n stands for network
- s stands for short      l stands for long

# bind ( )

---

- `int bind(int sockfd, struct sockaddr *my_addr,int addrlen);`
- `sockfd` - the socket file descriptor returned by `socket()`.
- `my_addr` - a pointer to a struct `sockaddr` that contains information about IP address and port number.
- `addrlen` - set to `sizeof (struct sockaddr)`

# connect( )

---

- `int connect (int sockfd, struct sockaddr *serv_addr, int addrlen);`
- `sockfd` - the socket file descriptor returned by `socket()`.
- `serv_addr` - is a `struct sockaddr` containing the destination port and IP address.
- `addrlen` - set to `sizeof (struct sockaddr)`.

# listen( )

---

- `int listen (int sockfd,int backlog);`
- sockfd - the socket file descriptor returned by socket().
- backlog - the number of connections allowed on the incoming queue.
- Backlog should never be zero as servers always expect connection from client.
- The listen function converts an unconnected socket into a passive socket,
- On successful execution of listen is indicating that the kernel should accept incoming connection requests directed to this socket.

# accept( )

---

- **int accept (int sockfd, void \*addr, int \*addrlen);**
- sockfd
  - the socket file descriptor returned by socket().
- addr
  - a pointer to a struct sockaddr\_in. The information about the incoming connection like IP address and port number are stored.
- addrlen
  - a local integer variable that should be set to sizeof (struct sockaddr\_in) before its address is passed to accept().

# close ( )

---

- Socket descriptor can be closed like file descriptor.
- `close (sockfd);`
- Close system call prevents any more reads and writes to the socket. For attempting to read or write the socket on the remote end will receive an error.



# shutdown ( )

---

- `int shutdown (int sockfd, int how);`
- sockfd - socket file descriptor of the socket to be shutdown.
- how – if it is
  - 0 - Further receives are disallowed
  - 1 - Further sends are disallowed
  - 2 - Further sends and receives are disallowed.
- The shutdown system call gives more control (than close (sockfd) over how the socket descriptor can be closed.

# pseudo code

## SERVER

```
struct sockaddr_in serv, cli;

sd = socket (AF_INET, SOCK_STREAM, 0);

serv.sin_family = AF_INET;
serv.sin_addr.s_addr = INADDR_ANY;
serv.sin_port = htons (portno);

bind (sd, &serv, sizeof (serv));
listen (sd, 5);

nsd = accept (sd, &cli, &sizeof (cli));
read / write (nsd,);
```

## CLIENT

```
struct sockaddr_in serv;

sd = socket (AF_INET, SOCK_STREAM, 0);

serv.sin_family = AF_INET;
serv.sin_addr.s_addr = inet_addr ("ser ip");
serv.sin_port = htons (portno);

connect (sd, &server, sizeof (server));

read / write (sd,);
```

# Iterative Server

---

- One client request at a time.

```
while (1) {
 nsd = accept (sd, &cli,...);
 read/write(nsd, ...);
}
```

# Concurrent Server

---

- Many clients requests can be serviced concurrently

```
while (1) {
 nsd =(accept (sd, &cli,));
 if (!fork()) {
 close(sd);
 read/write(nsd,);
 exit();
 } else
 close(nsd);
}
```

# Bibilography

---

1. Stevens. W R. Unix Network Programming Volume I & II
2. Chan, Terrance. Unix System Programming Using C++
3. Bach, Maurice J. Design of Unix

## **References :**

<http://linuxdocs.org/HOWTOs/DB2-HOWTO/kernel24.html>



**Thank You**